
title: 【LLM】Qwen3.5解剖 url:

<https://zhuanlan.zhihu.com/p/2005306558997882654> author: Plunck

author_badge: AI & LLMs created: 2026-02-13 03:22 modified: 2026-02-15
13:01 upvote_num: 245 comment_num: 11


目录

1. 一、文本模型

1. 1. MRoPE多模态（时间、高度、宽度）旋转位置编码
2. 2. 全注意力Qwen3_5Attention
3. 3. 线性注意力Qwen3_5GatedDeltaNet

2. 二、视觉模型

Figure 1: Qwen3.5 Overall Architecture



Qwen3.5 即将发布，从代码结构上看是Qwen 家族第一个原生多模态模型，后续应该不会再有 Qwen3.5-VL 了。对比前一代模型主要有以下几个创新点：


- **混合层架构**：注意力使用了 Qwen3-Next 里的混合层架构，每3层线性注意力插入1层标准注意力
- **Gated Attention**：全注意力和线性注意力都引入了Qwen3-Next 中的门控机制。其中线性注意力采用了DeltaNet
- **原生多模态**：大概率从预训练开始使用了视觉数据

[Code] https://github.com/bozheng-hit/transformers/blob/f3bd443cc9274ac8550b4585114323200b780f5a/src/transformers/models/qwen3_5/modeling_qwen3_5.py

之前的Qwen 系列解读参考：

[【多模态大模型】Qwen3-VL解剖](#)

Figure 1: Qwen3.5 Overall Architecture



模型概览

一、文本模型

Qwen3_5TextModel 混合了全注意力 (Full Attention) 和线性注意力 (Linear Attention / Gated DeltaNet)，并使用多模态旋转位置编码 MRoPE。在模型初始化阶段，构建了以下关键组件：

```
def __init__(self, config: Qwen3_5TextConfig):
    super().__init__(config)
    self.embed_tokens = nn.Embedding(config.vocab_size, config.hidden_size, config.pad_token_id)
```

```

self.layers = nn.ModuleList(
    [Qwen3_5DecoderLayer(config, layer_idx) for layer_idx in range(config.num_hidden_layers)]
)
self.norm = Qwen3_5RMSNorm(config.hidden_size, eps=config.rms_norm_eps)
self.rotary_emb = Qwen3_5TextRotaryEmbedding(config=config)
self.gradient_checkpointing = False
# Initialize weights and apply final processing
self.post_init()

```

1. MRoPE多模态（时间、高度、宽度）旋转位置编码

其中值得关注的是 `Qwen3_5TextRotaryEmbedding`，熟悉Qwen-VL系列模型的朋友知道这是一个支持多模态（时间、高度、宽度）的旋转位置编码。对于给定的 `position_ids`（纯文本 `[batch, seq_len]` 或多模态 `[3, batch, seq_len]`），它返回对应的cos和sin矩阵，用于与 `query/key` 向量进行旋转。初始化方法中定义了两个比较重要的变量：

- **partial_rotary_factor**：配置文件里写 0.25，意义是只对 O/K 四分之一的维度进行旋转，具体来说，注意力每个头的维度是 `head_dim=256`，用于旋转编码的维度是 `dim=256*0.25=64`。

```

inv_freq = 1.0 / (base ** (torch.arange(0, dim, 2, dtype=torch.int64).to(device=device, dtype=torch.float) / dim))

```

- **inv_freq**：逆频率，形状 `[D]`，其中 $D = \text{dim} // 2 = 32$ ，是旋转频率的个数，RoPE 的核心参数，控制不同维度正弦波形的周期。

接下来看 `Qwen3_5TextRotaryEmbedding` 的forward部分就比较好理解了：

```

def forward(self, x, position_ids):
    # In contrast to other models, Qwen3_5 has different position ids for the grids
    # So we expand the inv_freq to shape (3, ...)
    if position_ids.ndim == 2:
        position_ids = position_ids[None, ...].expand(3, position_ids.shape[0], -1)
    inv_freq_expanded = self.inv_freq[None, None, :, None].float().expand(3, position_ids.shape[1], -1, 1)
    position_ids_expanded = position_ids[:, :, None, :].float() # shape (3, bs, 1, positions)

    device_type = x.device.type if isinstance(x.device.type, str) and x.device.type != "mps" else "cpu"
    with maybe_autocast(device_type=device_type, enabled=False): # Force float32
        freqs = (inv_freq_expanded.float() @ position_ids_expanded.float()).transpose(2, 3)
        freqs = self.apply_interleaved_mrope(freqs, self.mrope_section)
        emb = torch.cat((freqs, freqs), dim=-1)
        cos = emb.cos() * self.attention_scaling
        sin = emb.sin() * self.attention_scaling

    return cos.to(dtype=x.dtype), sin.to(dtype=x.dtype)

```

其中expanded的维度变化如下：

- **inv_freq**：初始形状 `[D]`。经过四次索引：`[None, None, :, None]` → 形状 `[1, 1, D, 1]`。接着 `.expand(3, bs, -1, 1)` → 形状 `[3, bs, D, 1]`。
- **position_ids**：原始形状可能是 `[bs, seq_len]`（纯文本）或 `[3, bs, seq_len]`（多模态）。代码保证扩展为3维，经过 `[:, :, None, :]` → 形状 `[3, bs, 1, seq_len]`。

接下来是矩阵乘法：

```

freqs = (inv_freq_expanded.float() @ position_ids_expanded.float()).transpose(2, 3)

```

这一步得到freqs的中间结果，形状`[3, bs, D, seq_len]`，离我们的期望仍有差异，因为时间、高度、宽度散落在三个不同通道。

`apply_interleaved_mrope`可以将三个通道的旋转角度汇总到一个通道：

```
def apply_interleaved_mrope(self, freqs, mrope_section):
    """Apply interleaved MROPE to 3D rotary embeddings.
    Reorganizes frequency layout from chunked [TTT...HHH...WWW] to
    interleaved [THWTHWTHW...TT], preserving frequency continuity.
    args:
        x: (3, bs, seq_len, head_dim // 2)
        mrope_section: (3,)
    returns:
        x_t: (bs, seq_len, head_dim // 2)
    """
    freqs_t = freqs[0] # just overwrite the first dimension T
    for dim, offset in enumerate((1, 2), start=1): # H, W
        length = mrope_section[dim] * 3
        idx = slice(offset, length, 3)
        freqs_t[..., idx] = freqs[dim, ..., idx]
    return freqs_t
```

`mrope_section` 是一个长度为 3 的列表，表示在旋转维度中分配给 **T、H、W** 各自的数量。例如 `[11, 11, 10]` 表示总旋转维度为 $(11+11+10)=32$ 。

这个参数设计的原因是因为我们希望最终的频率矩阵（形状 `[bs, seq_len, D]`）中通道的顺序是：

```
[ T1, H1, W1, T2, H2, W2, T3, H3, W3, ..., T_{T_len}, H_{H_len}, W_{W_len} ]
```

每 3 个通道为一组，除非某维度先耗尽（这里是 W 少一个），同一组内的通道来自不同的模态。至此，`freqs` 的形状是我们想要的 `[bs, seq_len, D]` 了。

接下来是将每个频率对应的角度值复制一份变成 `[bs, seq_len, dim]` 并计算 $\cos/\sin$ ：

```
emb = torch.cat((freqs, freqs), dim=-1)
cos = emb.cos() * self.attention_scaling
sin = emb.sin() * self.attention_scaling
```

以便在后续的旋转计算中与向量拆分方式相匹配。这是常用的 RoPE 计算范式，这种范式要求 $\cos$ 和 $\sin$ 张量的维度与 `query` 完全一致，并且前半部分和后半部分的相同索引位置必须对应同一个频率 θ_i 。以 `query` 为例，`q_embed = q * cos + rotate_half(q) * sin`，其中 `rotate_half`：

```
def rotate_half(x):
    """Rotates half the hidden dims of the input."""
    x1 = x[..., : x.shape[-1] // 2]
    x2 = x[..., x.shape[-1] // 2 :]
    return torch.cat((-x2, x1), dim=-1)
```

代入上述 $\cos$ 和 $\sin$ ：

- $q * \cos$ ：前半 D 维乘 $\cos \theta_i$ ，后半 D 维也乘 $\cos \theta_i$ 。
- $\text{rotate_half}(q) * \sin$ ：前半 D 维 = $-q[\dots, D+i] * \sin \theta_i$ ，后半 D 维 = $q[\dots, i] * \sin \theta_i$ 。

合并后：

- 第 i 维 ($0 \leq i < D$)： $q_i * \cos \theta_i - q_{D+i} * \sin \theta_i$
- 第 $D+i$ 维： $q_{D+i} * \cos \theta_i + q_i * \sin \theta_i$

这正是将维度对 $(i, D+i)$ 作为旋转对，并用 θ_i 进行旋转的数学形式。

2. 全注意力Qwen3_5Attention

Qwen3_5DecoderLayer 每一层可能是linear_attention (GatedDeltaNet) 也可能是full_attention。具体来说每间隔四层有一个全注意力层，如 [L,L,L,F,L,L,L,F,L,...]。

```
class Qwen3_5DecoderLayer(GradientCheckpointingLayer):
    def __init__(self, config: Qwen3_5TextConfig, layer_idx: int):
        super().__init__()
        self.hidden_size = config.hidden_size
        self.layer_type = config.layer_types[layer_idx]
        if self.layer_type == "linear_attention":
            self.linear_attn = Qwen3_5GatedDeltaNet(config, layer_idx)
        elif self.layer_type == "full_attention":
            self.self_attn = Qwen3_5Attention(config, layer_idx)
        self.mlp = Qwen3_5MLP(config, config.intermediate_size)
        self.input_layernorm = Qwen3_5RMSNorm(config.hidden_size, eps=config.rms_norm_eps)
        self.post_attention_layernorm = Qwen3_5RMSNorm(config.hidden_size, eps=config.rms_norm_eps)
```

Hybrid Attention Mechanisms in Qwen3.5

Full Attention (Partial RoPE, QK-Norm, σ -Gating) & Gated DeltaNet Linear Attention




Figure 2: (a) Full Attention uses QK-Norm, Partial RoPE (25%=64d), GQA, and sigmoid output gating. (b) Gated DeltaNet uses causal Conv1D on joint QKV, L2-normalized Q/K, gated delta rule recurrence — $O(n)$ complexity.

首先讲一下全注意力** Qwen3_5Attention **的具体实现，这是一个普通的门控注意力。Gated Attention是NeurIPS 2025 Best Paper，之前做过详细讲解。

```
class Qwen3_5Attention(nn.Module):
    """Multi-headed attention from 'Attention Is All You Need' paper"""

    def __init__(self, config: Qwen3_5Config, layer_idx: int):
```

```

super().__init__()
self.config = config
self.layer_idx = layer_idx
self.head_dim = getattr(config, "head_dim", config.hidden_size // config.num_attention_head
s)

self.num_key_value_groups = config.num_attention_heads // config.num_key_value_heads
self.scaling = self.head_dim**-0.5
self.attention_dropout = config.attention_dropout
self.is_causal = True
self.q_proj = nn.Linear(
    config.hidden_size, config.num_attention_heads * self.head_dim * 2, bias=config.attentio
n_bias
)
self.k_proj = nn.Linear(
    config.hidden_size, config.num_key_value_heads * self.head_dim, bias=config.attention_bi
as
)
self.v_proj = nn.Linear(
    config.hidden_size, config.num_key_value_heads * self.head_dim, bias=config.attention_bi
as
)
self.o_proj = nn.Linear(
    config.num_attention_heads * self.head_dim, config.hidden_size, bias=config.attention_bi
as
)

self.q_norm = Qwen3_5RMSNorm(self.head_dim, eps=config.rms_norm_eps) # unlike olmo, only on
the head dim!
self.k_norm = Qwen3_5RMSNorm(self.head_dim, eps=config.rms_norm_eps) # thus post q_norm doe
s not need reshape

def forward(
    self,
    hidden_states: torch.Tensor,
    position_embeddings: tuple[torch.Tensor, torch.Tensor],
    attention_mask: torch.Tensor | None,
    past_key_values: Cache | None = None,
    cache_position: torch.LongTensor | None = None,
    **kwargs: Unpack[FlashAttentionKwargs],
) -> tuple[torch.Tensor, torch.Tensor | None]:
    input_shape = hidden_states.shape[:-1]
    hidden_shape = (*input_shape, -1, self.head_dim)

    query_states, gate = torch.chunk(
        self.q_proj(hidden_states).view(*input_shape, -1, self.head_dim * 2), 2, dim=-1
    )
    gate = gate.reshape(*input_shape, -1)

    query_states = self.q_norm(query_states.view(hidden_shape)).transpose(1, 2)
    key_states = self.k_norm(self.k_proj(hidden_states).view(hidden_shape)).transpose(1, 2)
    value_states = self.v_proj(hidden_states).view(hidden_shape).transpose(1, 2)

    cos, sin = position_embeddings
    query_states, key_states = apply_rotary_pos_emb(query_states, key_states, cos, sin)

    if past_key_values is not None:
        # sin and cos are specific to RoPE models; cache_position needed for the static cache

```

```

        cache_kwargs = {"sin": sin, "cos": cos, "cache_position": cache_position}
        key_states, value_states = past_key_values.update(key_states, value_states, self.layer_i
dx, cache_kwargs)

        attention_interface: Callable = ALL_ATTENTION_FUNCTIONS.get_interface(
            self.config._attn_implementation, eager_attention_forward
        )

        attn_output, attn_weights = attention_interface(
            self,
            query_states,
            key_states,
            value_states,
            attention_mask,
            dropout=0.0 if not self.training else self.attention_dropout,
            scaling=self.scaling,
            **kwargs,
        )

        attn_output = attn_output.reshape(*input_shape, -1).contiguous()
        attn_output = attn_output * torch.sigmoid(gate)

        attn_output = self.o_proj(attn_output)
        return attn_output, attn_weights

```

其核心是定义一个和 注意力输出 `attn_output` 一样形状 `[batch, seq_len, n_head*head_dim]` 的 `gate` , 然后逐元素相乘：

```
attn_output = attn_output * torch.sigmoid(gate)
```

`gate` 的值由当前的 `hidden_state`——也即当前的待输出 token 决定，这使得门控注意力可以通过当前的 token 的状态来调整注意力输出，这可能是它能提升训练效果的一个原因。

```
query_states, gate = torch.chunk(self.q_proj(hidden_states).view(*input_shape, -1, self.head_dim *
2), 2, dim=-1)
```

`apply_rotary_pos_emb` 对Query 和 Key 指定的维度（前四分之一）进行了旋转，其它的细节在上一节已详细讲过：

```
def apply_rotary_pos_emb(q, k, cos, sin, unsqueeze_dim=1):
    """Applies Rotary Position Embedding to the query and key tensors.

    Removes the interleaving of cos and sin from GLM

    Args:
        q (`torch.Tensor`): The query tensor.
        k (`torch.Tensor`): The key tensor.
        cos (`torch.Tensor`): The cosine part of the rotary embedding.
        sin (`torch.Tensor`): The sine part of the rotary embedding.
        unsqueeze_dim (`int`, *optional*, defaults to 1):
            The 'unsqueeze_dim' argument specifies the dimension along which to unsqueeze cos[position
            ids] and
            sin[position_ids] so that they can be properly broadcasted to the dimensions of q and k.
            For example, note
            that cos[position_ids] and sin[position_ids] have the shape [batch_size, seq_len, head_d
            im]. Then, if q and
            k have the shape [batch_size, heads, seq_len, head_dim], then setting unsqueeze_dim=1 ma

```



```

kes
        cos[position_ids] and sin[position_ids] broadcastable to the shapes of q and k. Similarly, if q and k have
        the shape [batch_size, seq_len, heads, head_dim], then set unsqueeze_dim=2.
    Returns:
        `tuple(torch.Tensor)` comprising of the query and key tensors rotated using the Rotary Position Embedding.
    """
    cos = cos.unsqueeze(unsqueeze_dim)
    sin = sin.unsqueeze(unsqueeze_dim)

    # Keep half or full tensor for later concatenation
    rotary_dim = cos.shape[-1]
    q_rot, q_pass = q[..., :rotary_dim], q[..., rotary_dim:]
    k_rot, k_pass = k[..., :rotary_dim], k[..., rotary_dim:]

    # Apply rotary embeddings on the first half or full tensor
    q_embed = (q_rot * cos) + (rotate_half(q_rot) * sin)
    k_embed = (k_rot * cos) + (rotate_half(k_rot) * sin)

    # Concatenate back to full shape
    q_embed = torch.cat([q_embed, q_pass], dim=-1)
    k_embed = torch.cat([k_embed, k_pass], dim=-1)
    return q_embed, k_embed

```

3. 线性注意力Qwen3_5GatedDeltaNet

Qwen3_5GatedDeltaNet 是 Qwen-3.5 引入的一种线性注意力（Linear Attention）或者说状态空间模型（SSM）机制。它的核心思想是使用 **Delta Rule（增量规则）** 来更新记忆状态，而不是像标准 Transformer 那样简单地累加（Add）。这种机制结合了 RNN 的推理效率 $O(1)$ 和 Transformer 的训练并行性（通过 Chunkwise 算法）。




Figure 2: An illustration of DeltaNet neural architecture.

3.1 初始化

首先看看初始化代码里有哪些内容，注释中补充了 config 里的部分超参：

```
class Qwen3_5GatedDeltaNet(nn.Module):
    def __init__(self, config: Qwen3_5Config, layer_idx: int):
        super().__init__()
        self.hidden_size = config.hidden_size # 4096
        self.num_v_heads = config.linear_num_value_heads # 32
        self.num_k_heads = config.linear_num_key_heads # 16
        # 这里允许 num_v_heads > num_k_heads, 通过 repeat_interleave 将 key 头广播到更多 value 头, 以提升容量而不显著增加计算量。
        self.head_k_dim = config.linear_key_head_dim # 128
        self.head_v_dim = config.linear_value_head_dim # 128
        self.key_dim = self.head_k_dim * self.num_k_heads # 2048
        self.value_dim = self.head_v_dim * self.num_v_heads # 4096

        self.conv_kernel_size = config.linear_conv_kernel_dim # 4
        self.layer_idx = layer_idx
        self.activation = config.hidden_act # "silu"
        self.act = ACT2FN[config.hidden_act]
        self.layer_norm_epsilon = config.rms_norm_eps

        # QKV
        self.conv_dim = self.key_dim * 2 + self.value_dim # 9192
        # 用矩阵转置使得一维卷积在序列维度 L 上进行
```

```

# (key_dim * 2 + value_dim) 作为通道数
self.conv1d = nn.Conv1d(
    in_channels=self.conv_dim,
    out_channels=self.conv_dim,
    bias=False,
    kernel_size=self.conv_kernel_size,
    groups=self.conv_dim,
    padding=self.conv_kernel_size - 1,
)

# time step projection (discretization)
# instantiate once and copy inv_dt in init_weights of PretrainedModel
self.dt_bias = nn.Parameter(torch.ones(self.num_v_heads))

A = torch.empty(self.num_v_heads).uniform_(0, 16)
self.A_log = nn.Parameter(torch.log(A))

self.norm = (
    Qwen3_5RMSNormGated(self.head_v_dim, eps=self.layer_norm_epsilon)
    if FusedRMSNormGated is None
    else FusedRMSNormGated(
        self.head_v_dim,
        eps=self.layer_norm_epsilon,
        activation=self.activation,
        device=torch.cuda.current_device(),
        dtype=config.dtype if config.dtype is not None else torch.get_default_dtype(),
    )
)

self.out_proj = nn.Linear(self.value_dim, self.hidden_size, bias=False)

self.causal_conv1d_fn = causal_conv1d_fn # 如有就从 fla 引入, 否则退化成上面定义的self.conv1d
self.causal_conv1d_update = causal_conv1d_update or torch.causal_conv1d_update
self.chunk_gated_delta_rule = chunk_gated_delta_rule or torch.chunk_gated_delta_rule
self.recurrent_gated_delta_rule = fused_recurrent_gated_delta_rule or torch.recurrent_gated_delta_rule

if not is_fast_path_available:
    logger.warning_once(
        "The fast path is not available because one of the required library is not installed. Falling back to "
        "torch implementation. To install follow https://github.com/fla-org/flash-linear-attention#installation and "
        "https://github.com/Dao-AILab/causal-conv1d"
    )

self.in_proj_qkv = nn.Linear(self.hidden_size, self.key_dim * 2 + self.value_dim, bias=False)
self.in_proj_z = nn.Linear(self.hidden_size, self.value_dim, bias=False)
self.in_proj_b = nn.Linear(self.hidden_size, self.num_v_heads, bias=False)
self.in_proj_a = nn.Linear(self.hidden_size, self.num_v_heads, bias=False)

```

`self.dt_bias` 和 `self.A_log` 是两个可学习的向量。

`Qwen3_5RMSNormGated` 是一个特殊的归一化层, 特殊在于将归一化后的 `attn_output` 与门控信号 `gate` 相乘, 位置和与全注意力的门控一样, 都在 `attention` 之后, `$W_O$` 之前。

```

class Qwen3_5RMSNormGated(nn.Module):
    def __init__(self, hidden_size, eps=1e-6, **kwargs):
        super().__init__()
        self.weight = nn.Parameter(torch.ones(hidden_size))
        self.variance_epsilon = eps

    def forward(self, hidden_states, gate=None):
        input_dtype = hidden_states.dtype
        hidden_states = hidden_states.to(torch.float32)
        variance = hidden_states.pow(2).mean(-1, keepdim=True)
        # Norm before gate
        hidden_states = hidden_states * torch.rsqrt(variance + self.variance_epsilon)
        hidden_states = self.weight * hidden_states.to(input_dtype)
        hidden_states = hidden_states * F.silu(gate.to(torch.float32))

        return hidden_states.to(input_dtype)

```

`torch_causal_conv1d_update` 完成一维卷积，并更新缓存的 `conv_state`：

```

def torch_causal_conv1d_update(
    hidden_states,
    conv_state,
    weight,
    bias=None,
    activation=None,
):
    _, hidden_size, seq_len = hidden_states.shape
    state_len = conv_state.shape[-1]

    # 1. 拼接历史与当前输入
    # 将缓存的过去几个 Token (conv_state) 和当前的新 Token (hidden_states) 拼在一起
    hidden_states_new = torch.cat([conv_state, hidden_states], dim=-1).to(weight.dtype)

    # 2. 更新缓存 (Rolling Cache)
    # 取拼接后序列的最后 state_len 个元素，覆盖回 conv state
    conv_state.copy_(hidden_states_new[:, :, -state_len:])

    # 3. 执行卷积
    # groups=hidden_size 表示这是 Depthwise Convolution (深度卷积)
    # 每个通道独立进行卷积，互不干扰，计算量极小
    out = F.conv1d(hidden_states_new, weight.unsqueeze(1), bias, padding=0, groups=hidden_size)

    # 取最后 seq_len 个结果作为输出
    out = F.silu(out[:, :, -seq_len:])
    out = out.to(hidden_states.dtype)
    return out

```

3.2 前向传播

```

def forward(
    self,
    hidden_states: torch.Tensor,
    cache_params: Qwen3_5DynamicCache | None = None,
    cache_position: torch.LongTensor | None = None,
    attention_mask: torch.Tensor | None = None,
):

```

```

hidden_states = apply_mask_to_padding_states(hidden_states, attention_mask)

# Set up dimensions for reshapes later
batch_size, seq_len, _ = hidden_states.shape

use_precomputed_states = (
    cache_params is not None
    and cache_params.has_previous_state
    and seq_len == 1
    and cache_position is not None
)

# getting projected states from cache if it exists
if cache_params is not None:
    conv_state = cache_params.conv_states[self.layer_idx]
    recurrent_state = cache_params.recurrent_states[self.layer_idx]

mixed_qkv = self.in_proj_qkv(hidden_states)
mixed_qkv = mixed_qkv.transpose(1, 2)

z = self.in_proj_z(hidden_states)
z = z.reshape(batch_size, seq_len, -1, self.head_v_dim)

b = self.in_proj_b(hidden_states)
a = self.in_proj_a(hidden_states)

if use_precomputed_states:
    # 2. Convolution sequence transformation
    # NOTE: the conv state is updated in `causal_conv1d_update`
    mixed_qkv = self.causal_conv1d_update(
        mixed_qkv,
        conv_state,
        self.conv1d.weight.squeeze(1),
        self.conv1d.bias,
        self.activation,
    )
else:
    if cache_params is not None:
        conv_state = F.pad(mixed_qkv, (self.conv_kernel_size - mixed_qkv.shape[-1], 0))
        cache_params.conv_states[self.layer_idx] = conv_state
    if self.causal_conv1d_fn is not None:
        mixed_qkv = self.causal_conv1d_fn(
            x=mixed_qkv,
            weight=self.conv1d.weight.squeeze(1),
            bias=self.conv1d.bias,
            activation=self.activation,
            seq_idx=None,
        )
    else:
        mixed_qkv = F.silu(self.conv1d(mixed_qkv)[: , : , :seq_len])

mixed_qkv = mixed_qkv.transpose(1, 2)
query, key, value = torch.split(
    mixed_qkv,
    [
        self.key_dim,

```

```

        self.key_dim,
        self.value_dim,
    ],
    dim=-1,
)

query = query.reshape(batch_size, seq_len, -1, self.head_k_dim)
key = key.reshape(batch_size, seq_len, -1, self.head_k_dim)
value = value.reshape(batch_size, seq_len, -1, self.head_v_dim)

beta = b.sigmoid()
# If the model is loaded in fp16, without the .float() here, A might be -inf
q = -self.A.log().float().exp() * F.softplus(a.float() + self.dt_bias)
if self.num_v_heads // self.num_k_heads > 1:
    query = query.repeat_interleave(self.num_v_heads // self.num_k_heads, dim=2)
    key = key.repeat_interleave(self.num_v_heads // self.num_k_heads, dim=2)

if not use_precomputed_states:
    core_attn_out, last_recurrent_state = self.chunk_gated_delta_rule(
        query,
        key,
        value,
        g=g,
        beta=beta,
        initial_state=None,
        output_final_state=cache_params is not None,
        use_qk_l2norm_in_kernel=True,
    )
else:
    core_attn_out, last_recurrent_state = self.recurrent_gated_delta_rule(
        query,
        key,
        value,
        g=g,
        beta=beta,
        initial_state=recurrent_state,
        output_final_state=cache_params is not None,
        use_qk_l2norm_in_kernel=True,
    )

# Update cache
if cache_params is not None:
    cache_params.recurrent_states[self.layer_idx] = last_recurrent_state

# reshape input data into 2D tensor
core_attn_out = core_attn_out.reshape(-1, self.head_v_dim)
z = z.reshape(-1, self.head_v_dim)
core_attn_out = self.norm(core_attn_out, z)
core_attn_out = core_attn_out.reshape(batch_size, seq_len, -1)

output = self.out_proj(core_attn_out)
return output

```

该模块的大致数据流向如下：

1. 输入投影：将输入 x 映射为 q, k, v 以及门控信号 z (输出门的 Gate)、 b (用于计算 β) 和 a (用于计算衰减率)。
2. 短卷积 (Short Conv)：在 q, k, v 序列上应用步长为 4 的 1D 卷积，捕捉局部特征 (类似 Mamba)。
3. 参数计算：计算步长/写入强度 β_t 和衰减率 g_t 。
4. 核心 Delta Rule：更新隐藏状态 S 并计算注意力输出。
5. 输出门控与归一化：使用 RMSNorm 和门控 z 处理输出。

理解 Qwen3_5GatedDeltaNet 的核心就在于理解它如何通过两套不同的算法路径，分别处理 **并行训练/预填充 (Parallel)** 和 **串行推理 (Recurrent)**。

1、解码阶段 (Decoding Phase)

场景：模型已经处理完 Prompt，现在开始逐个生成 Token ($L+1, L+2, \dots$)。

数据特征：

- 输入形状：[Batch, 1, Hidden_Size] (每次只进 1 个 Token)。
- use_precomputed_states 为 True (前提是 cache_params 存在且非空)。

代码与数据流向：

- 读取 Cache：

```
if cache_params is not None:
    conv_state = cache_params.conv_states[self.layer_idx] # 形状 [B, D, Kernel]
    recurrent_state = cache_params.recurrent_states[self.layer_idx] # 形状 [B, H, K_dim, V_dim]
```

- 投影 (Projections)：对这 1 个 Token 进行投影，得到 $q_t, k_t, v_t, z_t, b_t, a_t$ 。
- 卷积 (Convolution) - 增量模式**：进入 if use_precomputed_states: 分支，调用 self.causal_conv1d_update。如前所述，<1>取出 conv_state (过去几个 Token)。<2>拼接当前 Token。<3>做卷积计算当前结果。<4>**原地更新** conv_state (扔掉最旧的，加入当前的)。在不重新计算整个序列的情况下，获得当前 Token 的卷积特征。

SSM 核心计算 - Recurrent 算法：

- 代码：调用 self.recurrent_gated_delta_rule。
- 逻辑：使用 RNN 形式的 Delta Rule 公式。
- $S_t = \lambda_t S_{t-1} + \beta_t k_t (v_t - S_{t-1})^T k_t^T$
- $o_t = S_t^T q_t$
- 输入：当前的 q, k, v 和上一时刻的 recurrent_state (S_{t-1})。
- 输出：core_attn_out (当前 Token 的输出 o_t) 和 last_recurrent_state (更新后的状态 S_t)。
- Cache 更新：cache_params.recurrent_states[self.layer_idx] = last_recurrent_state。

解码阶段完全是 $O(1)$ 的。无论之前的序列有多长 (1k 还是 100k)，计算量都只和 Hidden Size 有关，显存占用也是恒定的 (只存 S 矩阵和极小的 Conv State)。

2、预填充阶段 (Prefilling Phase)

场景：用户输入了一段 Prompt (例如长度 $L=1024$)，模型需要一次性计算出这段 Prompt 的所有 Token 的特征，并生成初始的记忆状态 (Cache)，为后续生成做准备。

数据特征：

- 输入形状：[Batch, Seq Len, Hidden_Size]，其中 Seq_Len > 1。
- use_precomputed_states 为 False。

投影 (Projections)：

- 输入 hidden_states 经过线性层，并行计算出整个序列的 Q, K, V 以及门控信号 z, b, a
- mixed_qkv = mixed_qkv.transpose(1, 2) 此时数据形状变为 [B, D, L]。

卷积 (Convolution) - 并行模式：

- **Cache 初始化**：用 F.pad 同时处理“序列过短需要补零”和“序列过长需要截取”两种情况，保持 conv_state 长度为 conv_kernel_size=4

```

if cache_params is not None:
    conv_state = F.pad(mixed_qkv, (self.conv_kernel_size - mixed_qkv.shape[-1], 0))
    cache_params.conv_states[self.layer_idx] = conv_state

```

- **操作**：调用 `self.causal_conv1d_fn` (如果安装了 CUDA 优化版) 或 `self.conv1d` (PyTorch 原生)。
- **逻辑**：在整个序列长度 L 上直接进行 1D 卷积。

```

else:
    mixed_qkv = F.silu(self.conv1d(mixed_qkv)[: , : , :seq_len])

```

参数计算：

- 恢复 q, k, v 的形状

```

mixed_qkv = mixed_qkv.transpose(1, 2)
query, key, value = torch.split(
    mixed_qkv,
    [
        self.key_dim,
        self.key_dim,
        self.value_dim,
    ],
    dim=-1,
)

query = query.reshape(batch_size, seq_len, -1, self.head_k_dim)
key = key.reshape(batch_size, seq_len, -1, self.head_k_dim)
value = value.reshape(batch_size, seq_len, -1, self.head_v_dim)

```

- 计算整个序列的 β (步长) 和 g (衰减率)。

```

beta = b.sigmoid()
# If the model is loaded in fp16, without the .float() here, A might be -inf
g = -self.A_log.float().exp() * F.softplus(a.float() + self.dt_bias)

```

- 在 DeltaNet 里， V 的头数比 Q 和 K 多，所以需要复制 Q/K 。

```

if self.num_v_heads // self.num_k_heads > 1:
    query = query.repeat_interleave(self.num_v_heads // self.num_k_heads, dim=2)
    key = key.repeat_interleave(self.num_v_heads // self.num_k_heads, dim=2)

```

V 的头数更多的原因是： **V 决定“写/读什么内容”（内容子空间容量）**，可以用更多的 value heads 提升表达能力/通道数，同时 g, β 是 per-value-head 的门控参数：让不同 value heads 即使共享 q, k ，也能用不同的衰减/更新强度形成差异化记忆。

SSM 核心计算 - Chunkwise 算法：

调用 `self.chunk_gated_delta_rule`。这是**并行化**的关键。它将长度 L 的序列切分成多个小块（Chunk=64）。**块内**使用类似 Attention 的矩阵乘法并行计算，**块间**传递隐藏状态。

```

def torch_chunk_gated_delta_rule(
    query,
    key,
    value,
    g,
    beta,
    chunk_size=64,
    initial_state=None,

```



```

output_final_state=False,
use_qk_l2norm_in_kernel=False,
):
    """
    Args:
        query, key, value: [Batch, Seq_Len, Num_Heads, Head_Dim]
        g: Log decay (衰减因子的对数), [Batch, Seq_Len, Num_Heads]
        beta: Step size (写入步长), [Batch, Seq_Len, Num_Heads]
        chunk_size: 分块大小, 通常为 64 或 128
    """
    initial_dtype = query.dtype

    # =====
    # 1. 预处理与归一化 (Preprocessing)
    # =====
    # 对 Q 和 K 进行 L2 归一化, 防止数值爆炸, 这是 Linear Attention 的常见技巧
    if use_qk_l2norm_in_kernel:
        query = l2norm(query, dim=-1, eps=1e-6)
        key = l2norm(key, dim=-1, eps=1e-6)

    # 转置为 [Batch, Num_Heads, Seq_Len, Head_Dim] 以便后续处理
    query, key, value, beta, g = [
        x.transpose(1, 2).contiguous().to(torch.float32) for x in (query, key, value, beta, g)
    ]

    batch_size, num_heads, sequence_length, k_head_dim = key.shape
    v_head_dim = value.shape[-1]

    # =====
    # 2. 填充 (Padding)
    # =====
    # 为了能整除 chunk_size, 对序列末尾进行 Padding
    pad_size = (chunk_size - sequence_length % chunk_size) % chunk_size
    query = F.pad(query, (0, 0, 0, pad_size))
    key = F.pad(key, (0, 0, 0, pad_size))
    value = F.pad(value, (0, 0, 0, pad_size))
    beta = F.pad(beta, (0, pad_size))
    g = F.pad(g, (0, pad_size))

    total_sequence_length = sequence_length + pad_size

    # Q 的缩放系数 (Scale), 类似 Attention 中的 1/sqrt(d)
    scale = 1 / (query.shape[-1] ** 0.5)
    query = query * scale

    # 预先计算加权的 V 和 K
    # v_beta = v * beta
    # k_beta = k * beta
    v_beta = value * beta.unsqueeze(-1)
    k_beta = key * beta.unsqueeze(-1)

    # =====
    # 3. 重塑为 Chunks (Reshape to Chunks)
    # =====
    # 将序列维度拆分为 [Num_Chunks, Chunk_Size]
    # 形状变为: [Batch, Heads, Num_Chunks, Chunk_Size, Dim]

```

```

query, key, value, k_beta, v_beta = [
    x.reshape(x.shape[0], x.shape[1], -1, chunk_size, x.shape[-1]) for x in (query, key, value,
k_beta, v_beta)
]
g = g.reshape(g.shape[0], g.shape[1], -1, chunk_size)

# 生成上三角掩码，用于后续的因果遮蔽
mask = torch.triu(torch.ones(chunk_size, chunk_size, dtype=torch.bool, device=query.device), diagonal=0)

# =====
# 4. 块内衰减计算 (Intra-Chunk Decay)
# =====
# 计算块内的累积衰减。
# g 是 log space, cumsum 后相减再 exp 得到相对衰减系数
g = g.cumsum(dim=-1)
# decay_mask[i, j] = exp(g[i] - g[j]), 表示从时刻 j 到 i 的衰减
decay_mask = ((g.unsqueeze(-1) - g.unsqueeze(-2)).tril().exp().float()).tril()

# =====
# 5. 块内正交化/逆变换 (Intra-Chunk Orthogonalization)
# 这是 Delta Rule 并行化的核心难点。
# =====
# Delta Rule 公式:  $v_{new} = v - (S_{t-1} * k)$ 
# 在块内并行时，我们需要解出“如果不考虑前一个块的状态，当前块产生的纯净更新量是多少”。
# 这等价于求解矩阵方程:  $(I + LowerTriangular) * V_{new} = V_{input}$ 

# 初始化 attn = - (K * beta * K^T) * Decay
attn = -((k_beta @ key.transpose(-1, -2)) * decay_mask).masked_fill(mask, 0)

# 通过迭代求解逆矩阵  $(I - A)^{-1} \approx I + A + A^2 \dots$ 
# 这里利用了下三角矩阵的性质，在 chunk_size 次内完成传播
for i in range(1, chunk_size):
    row = attn[..., i, :i].clone()
    sub = attn[..., :i, :i].clone()
    attn[..., i, :i] = row + (row.unsqueeze(-1) * sub).sum(-2)

# 加上单位矩阵 I
attn = attn + torch.eye(chunk_size, dtype=attn.dtype, device=attn.device)

# 计算块内的“净增量” (Net Delta)
# value 现在存储的是经过块内相互抵消后的 v_new
value = attn @ v_beta

# 预计算 K 的累积效应，用于计算对下一个块的状态贡献
#  $K_{cumdecay} = (I + \dots)^{-1} * K * beta * decay$ 
k_cumdecay = attn @ (k_beta * g.exp().unsqueeze(-1))

# =====
# 6. 块间递归 (Inter-Chunk Recurrence)
# 主循环：串行处理每个 Chunk，更新状态 S
# =====
# 初始化记忆状态 S (Hidden State)
last_recurrent_state = (
    torch.zeros(batch_size, num_heads, k_head_dim, v_head_dim).to(value)
    if initial_state is None

```

```

        else initial_state.to(value)
    )
    core_attn_out = torch.zeros_like(value)
    # 用于 Attention 的掩码 (对角线偏移1)
    mask = torch.triu(torch.ones(chunk_size, chunk_size, dtype=torch.bool, device=query.device), diagonal=1)

    # 遍历每一个 Chunk
    for i in range(0, total_sequence_length // chunk_size):
        # 取出当前 Chunk 的数据
        q_i, k_i, v_i = query[:, :, i], key[:, :, i], value[:, :, i]

        # --- A. 计算块内 Attention ---
        # 标准的 Q * K^T * Decay
        attn = (q_i @ k_i.transpose(-1, -2) * decay_mask[:, :, i]).masked_fill_(mask, 0)

        # --- B. 引入历史状态的影响 ---
        # 计算历史状态 S_{n-1} 对当前块 Key 的投影 (即预测值)
        # v_prime = S_{n-1}^T * K_cumdecay
        v_prime = (k_cumdecay[:, :, i]) @ last_recurrent_state

        # 计算最终的 Delta (误差)
        # v_new = 块内净增量 - 历史预测
        # 对应公式: delta = beta * (v - S * k)
        v_new = v_i - v_prime

        # --- C. 计算输出 (Output) ---
        # 第一部分: Q 读取历史状态 S_{n-1}
        # attn_inter = Q * Decay * S_{n-1}
        attn_inter = (q_i * g[:, :, i, :, None].exp()) @ last_recurrent_state

        # 第二部分: Q 读取当前块内的更新 v_new
        # out = Q * S_{n-1} + (Q * K^T) * v_new
        core_attn_out[:, :, i] = attn_inter + attn @ v_new

        # --- D. 更新状态 (State Update) ---
        # S_n = S_{n-1} * decay + K^T * v_new
        # 第一项: 历史状态衰减
        # 第二项: 写入新的信息 (外积 K^T * v_new)
        last_recurrent_state = (
            last_recurrent_state * g[:, :, i, -1, None, None].exp()
            + (k_i * (g[:, :, i, -1, None] - g[:, :, i]).exp()[..., None]).transpose(-1, -2) @ v_new
        )

    # =====
    # 7. 后处理 (Post-processing)
    # =====
    if not output_final_state:
        last_recurrent_state = None

    # 还原形状 [Batch, Heads, Num_Chunks, Chunk_Size, Dim] -> [Batch, Heads, Seq_Len, Dim]
    core_attn_out = core_attn_out.reshape(core_attn_out.shape[0], core_attn_out.shape[1], -1, core_attn_out.shape[-1])

    # 去除 Padding
    core_attn_out = core_attn_out[:, :, :sequence_length]

```

```
# 转置回 [Batch, Seq_Len, Heads, Dim]
core_attn_out = core_attn_out.transpose(1, 2).contiguous().to(initial_dtype)

return core_attn_out, last_recurrent_state
```

- **输出** : `core_attn_out` (整个序列的注意力输出 `[B, L, D]`) 和 `last_recurrent_state` (处理完第 `L` 个 token 后的最终记忆矩阵 `$S_L$`)。
- **Cache 保存** : 将最终状态 `S_L` 存入 `recurrent_states`。

预填充阶段利用 GPU 的并行能力，通过 `Chunkwise` 算法快速处理长序列，并产出“压缩”后的记忆状态。

最后，`prefilling` 和 `decoding` 阶段结束之后都需要对结果进行门控归一化，`$W_O$` 投影输出。

```
# reshape input data into 2D tensor
core_attn_out = core_attn_out.reshape(-1, self.head_v_dim)
z = z.reshape(-1, self.head_v_dim)
core_attn_out = self.norm(core_attn_out, z)
core_attn_out = core_attn_out.reshape(batch_size, seq_len, -1)

output = self.out_proj(core_attn_out)
return output
```

二、视觉模型

看了很久貌似只是去除了 `DeepStack`，其它和 `Qwen3-VL` 的视觉部分没什么区别。

模型层面的创新是比较稀疏的，因为不如 `scaling` 数据、搭建更真实的训练环境来得立竿见影，而且有些 `trick` 并不一定总是有效，有时候还面临尴尬的回退，比如这份代码里 `MRoPE` 依然是退化的 (`t_index is always 0`)，不知道为什么没被删掉。